



Concurrent programming in Rust

4SE05

Léopold Clément

Télécom Paris

2025-05-19

Outline

- Concurrency in Rust 2
 - Safty guarantees in Rust 5
- Concurrent primitives for threaded concurrency 17
 - Threads 18
 - Mutex 25
 - Channel `mpsc` 32
 - Usefull crates 38
- Async/await, Future and cooperative concurrence 41
 - Getting two pages at the same time 42
 - Implementation in Rust 49
- Conclusion 64

Concurrency in Rust



For example, introducing parallelism in Rust is a relatively low-risk operation: the compiler will catch the classical mistakes for you.

— The rust book, foreword

By leveraging ownership and type checking, many concurrency errors are compile-time errors in Rust rather than runtime errors.

— The rust book, chapter 16

Two types of concurrency

Two types of concurrency

- Threaded concurrency
 - Based on threads, like in C
 - Good for compute-bound workloads,

Two types of concurrency

- Threaded concurrency
 - Based on threads, like in C
 - Good for compute-bound workloads,
- Asynchronous concurrency
 - Based on future / promise / task / async functions
 - Seen in Python, JavaScript/TypeScript
 - Good for IO-bound workloads,
 - Used in embassy (cf 4SE02)

Concurrency in Rust

Safety guarantees in Rust

The borrow checker

The borrow checker guarantees that there is either:

- only one mutable reference to a value
- multiple immutables.

The borrow checker

The borrow checker guarantees that there is either:

- only one mutable reference to a value
- multiple immutables.

Raw object cannot be shared between threads and mutable.

The borrow checker guarantees that there is either:

- only one mutable reference to a value
- multiple immutables.

Raw object cannot be shared between threads and mutable.

To share object between threads and have mutability, you need to use some synchronisation primitive.

The borrow checker guarantees that there is either:

- only one mutable reference to a value
- multiple immutables.

Raw object cannot be shared between threads and mutable.

To share object between threads and have mutability, you need to use some synchronisation primitive.

Rc is a smart pointer.

It can be cloned, and counts the number of time it was cloned.

It allows for a value to have a lifetime long enough for each time it is used.

```
fn rc_to_thread() {  
    let a = Rc::from(1);  
    let b = a.clone();  
    let j = spawn(move || println!("val = {}", *b));  
    j.join();  
}
```

Sharing value with Rc, errors!

```
error[E0277]: `Rc<i32>` cannot be sent between threads safely
```

```
--> src/main.rs:90:19
```

```
90 |         let j = spawn(move || println!("val = {}", *b));
```

```
-----^
```

```
    |  
    |  
    | `Rc<i32>` cannot be sent between threads safely  
    | within this `{closure@src/main.rs:90:19: 90:26}`  
    | required by a bound introduced by this call
```

```
= help: within `{closure@src/main.rs:90:19: 90:26}`, the trait `Send`  
is not implemented for `Rc<i32>`
```

The Send and Sync traits

The compiler implements those traits when possible :

The Send and Sync traits

The compiler implements those traits when possible :

Send The value can be moved between threads.

The Send and Sync traits

The compiler implements those traits when possible :

Send The value can be moved between threads.

Sync Reference to the value can be moved between threads.

To share the ownership of a value of type `T` across thread boundaries, you might use `Arc<T>` instead of `Rc<T>`.

To share the ownership of a value of type `T` across thread boundaries, you might use `Arc<T>` instead of `Rc<T>`.

Here the counter is atomic, so the compiler adds the `Send` trait and the `Arc` can be moved between threads.

To share the ownership of a value of type `T` across thread boundaries, you might use `Arc<T>` instead of `Rc<T>`.

Here the counter is atomic, so the compiler adds the `Send` trait and the `Arc` can be moved between threads.

Shared ownership does not permit mutability, `Arc` can only be deref into immutable reference.

Most of the primitives are defined in the `std` module.

Most of the primitives are defined in the `std` module.

- `std`
 - `thread`
 - `sync`
 - `atomics`
 - `mutex`
 - `channel`
 - `future`
 - `task`

Most of the primitives are defined in the `std` module.

- `std`
 - `thread`
 - `sync`
 - `atomics`
 - `mutex`
 - `channel`
 - `future`
 - `task`

You can build more complex primitives from basic primitives, or using **unsafe** blocks.

Atomics live in `std::sync::atomic`.

```
use std::sync::atomic::{AtomicU16, Ordering};

let atomic = AtomicU16::new(0);
assert_eq!(atomic.load(Ordering::Relaxed), 0);
atomic.store(1, Ordering::Relaxed);
assert_eq!(atomic.fetch_add(10, Ordering::Relaxed), 1);
assert_eq!(atomic.swap(100, Ordering::Relaxed), 11);
assert_eq!(atomic.load(Ordering::Relaxed), 100);
```


Atomics live in `std::sync::atomic`.

```
use std::sync::atomic::{AtomicU16, Ordering};

let atomic = AtomicU16::new(0);
assert_eq!(atomic.load(Ordering::Relaxed), 0);
atomic.store(1, Ordering::Relaxed);
assert_eq!(atomic.fetch_add(10, Ordering::Relaxed), 1);
assert_eq!(atomic.swap(100, Ordering::Relaxed), 11);
assert_eq!(atomic.load(Ordering::Relaxed), 100);
```

The ordering is the same as in C++ 20

Relaxed: only guarantee atomicity and modification order consistency,
AcqRel: also, things that happen before/after a store in one thread happen before/after a read in another thread,
SeqCst: also, there is a **single total modification order** of all atomic operations that are so tagged

Atoms : Creating a mutex

```
pub struct SpinLockMutex{
    locked : AtomicBool,
}
impl SpinLockMutex {
    pub fn new() -> Self{
        SpinLockMutex { locked: AtomicBool::new(false) }
    }
    pub fn lock(&self) -> (){
        loop {
            let old_state = AtomicBool::swap(&self.locked, true, Ordering::AcqRel);
            if !old_state {return;}
        }
    }
    pub fn unlock(&self) -> (){
        AtomicBool::store(&self.locked, false, Ordering::AcqRel);
    }
}
```

```
let atomic = AtomicU16::new(0);
```

```
pub fn swap(&self, val: bool, order: Ordering) -> bool
```

```
pub fn lock(&self) -> ()
```

```
pub fn unlock(&self) -> ()
```

```
let atomic = AtomicU16::new(0);
```

```
pub fn swap(&self, val: bool, order: Ordering) -> bool
```

```
pub fn lock(&self) -> ()
```

```
pub fn unlock(&self) -> ()
```

In Rust, some types allow their values to be mutated even on immutable instances.

`Cell<T>` is a container used to create interior mutability.

```
pub fn set(&self, val: T)
pub fn replace(&self, val: T) -> T
pub fn swap(&self, other: &Cell<T>)
```

`Cell` is not `Send` : it is not atomic w.r.t other threads, only w.r.t its own thread.

Concurrent primitives for threaded concurrency



Concurrent primitives for threaded concurrency

Threads

Rust threads are os threads.

Rust threads are os threads.

- If the main thread panics, the program stops.

Rust threads are os threads.

- If the main thread panics, the program stops.
- If another thread panics, nothing happens.

Rust threads are os threads.

- If the main thread panics, the program stops.
- If another thread panics, nothing happens.
- If the main thread terminates before the other, the other threads are stopped.

FnOnce Can be called once.

FnOnce Can be called once.

Fn Can be called multiple time.

FnOnce Can be called once.

Fn Can be called multiple time.

FnMut Can be called multiple time and mutate borrowed value.

FnOnce Can be called once.

Fn Can be called multiple time.

FnMut Can be called multiple time and mutate borrowed value.

```
let mut x = 0;  
let add = |y| {x + y}; //Fn  
let mut inc_by = |y| {x += y}; //FnMut
```

```
x: String = String::from("Léopold");  
let concat_to_x = |y: String| add(x, &y); //FnOnce  
// x is consumed by add
```


Spawning and joining a thread

```
let t = std::thread::spawn(move || {  
    println!("Hi from the thread");  
});  
println!("Hi from main");  
t.join().unwrap();
```

Spawning and joining a thread

```
let t = std::thread::spawn(move || {  
    println!("Hi from the thread");  
});  
println!("Hi from main");  
t.join().unwrap();
```

Always join a thread.

Move value inside threads, wrong

```
let name = Arc::new("Leopold");  
let hello_from_inside = |x:usize, name| {println!("Hi {} from {}", name,  
x)};};  
let t1 = spawn(move || hello_from_inside(1, name));  
let t2 = spawn(move || hello_from_inside(2, name));  
error[E0382]: use of moved value: `name`
```

Arc must be explicitly cloned.

```
let name = Arc::new("Leopold");  
let hello_from_inside = |x:usize, name| {println!("Hi {} from {}",  
name, x)};};  
let tmp = name.clone(); // clone the Arc  
let t1 = spawn(move || hello_from_inside(1, tmp)); //tmp move, not  
name  
let tmp = name.clone();  
let t2 = spawn(move || hello_from_inside(2, tmp));
```

There are two kinds of communication:

There are two kinds of communication:

- Shared state
 - ▶ use mutex
 - ▶ every threads can read/write

There are two kinds of communication:

- Shared state
 - use `mutex`
 - every threads can read/write
- Message passing
 - use `channel`
 - some threads are writers, some are readers

Concurrent primitives for threaded concurrency

Mutex

A `mutex` allows to share a value between threads. Only one thread may hold a reference to the value at once.

A `mutex` allows to share a value between threads. Only one thread may hold a reference to the value at once. In Rust, `Mutex` is a wrapper around a non-atomic value. Each thread has a reference to the `Mutex`, but only one has a reference to the value.

Lifecycle of a Mutex

```
let mut threads : Vec<JoinHandle<_>> = Vec::new();
let counter = Arc::new(Mutex::new(0));
for idx_thread in 0..N_MAX{
    let counter = counter.clone();
    threads.push(spawn(move || {
        for _ in 0..idx_thread{
            let mut c = counter.lock().unwrap();
            *c += 1;}
        })));
}
for t in threads.into_iter(){t.join().unwrap();}
assert_eq!(*counter.lock().unwrap(), (N_MAX-1)*(N_MAX)/2);
```

Taking a Mutex

```
fn lock(&self) -> LockResult<MutexGuard<'_, T>>;  
fn try_lock(&self) -> TryLockResult<MutexGuard<'_, T>>;
```

```
LockResult<T> = Result<T, PoisonError<T>>;  
TryLockResult<T> = Result<T, TryLockError<T>>;
```

The mutex might be poisoned, and can not be taken.

Taking a Mutex

```
fn lock(&self) -> LockResult<MutexGuard<'_, T>>;  
fn try_lock(&self) -> TryLockResult<MutexGuard<'_, T>>;
```

```
LockResult<T> = Result<T, PoisonError<T>>;  
TryLockResult<T> = Result<T, TryLockError<T>>;
```

The mutex might be poisoned, and can not be taken.

There are two interfaces :

- blocking : `lock`
- non-blocking : `try_lock`

The mutex is **released at the end of the context** where it was locked.

You can also use `std::mem::drop`.

The mutex is **released at the end of the context** where it was locked.

You can also use `std::mem::drop`.

```
for _ in 0..idx_thread{  
    let mut c = counter.lock().unwrap();  
    *c += 1;  
} // counter is released here, on each iteration of the loop
```

If a thread panics while holding a mutex, the mutex will be poisoned. A poisoned mutex **cannot be taken** by anyone.

A mutex can be tested for poison with `is_poisoned` and cured with `clear_poison`. This only affects the mutex, the inner data might still be corrupted.

The `Mutex` allows for only one reference to the value at once.

Concurrent primitives for threaded concurrency

Channel `mpsc`

A `Channel` allows to pass messages between contexts.

The base implementation allows multiple writers and one reader. The channel has an infinite capacity.

Creating a channel

A channel for values of type `T` is represented by its ends, the `Sender<T>` (`tx`) and the `Receiver<T>` (`rx`).

```
let (tx, rx) = std::sync::mpsc::channel();
```

The `Sender` can be cloned, to create multiple producers.

Sending through the channel

```
fn send(&self, t: T) -> Result<(), SendError<T>>
```

Returns an error only if the `rx` is dropped. `Ok` does not mean that the message is received. This will never block.

Receiving through the channel

```
fn recv(&self) -> Result<T, RecvError>
```

Will block if no message in the channel. Will return an `Err` if the Sender is disconnected.

Receiving through the channel

```
fn recv(&self) -> Result<T, RecvError>
```

Will block if no message in the channel. Will return an Err if the Sender is disconnected.

There is a non-clocking interface, try_recv and a timeout interface recv_timeout.

```
fn recv_timeout(&self, timeout: Duration) -> Result<T, RecvTimeoutError>  
fn try_recv(&self) -> Result<T, TryRecvError>
```

Receiving through the channel

```
fn recv(&self) -> Result<T, RecvError>
```

Will block if no message in the channel. Will return an Err if the Sender is disconnected.

There is a non-clocking interface, try_recv and a timeout interface recv_timeout.

```
fn recv_timeout(&self, timeout: Duration) -> Result<T, RecvTimeoutError>  
fn try_recv(&self) -> Result<T, TryRecvError>
```

The Receiver can also be turned into an iterator:

```
for msg in rx.iter() {  
    ...  
}
```


A channel allows to move values from one thread to an other.

Concurrent primitives for threaded concurrency

Usefull crates

Another implementation of Mutex

- smaller
- fair

Also see the `RwLock`, for multiple immutable references or one mutable reference at the same time.

For data parallelism, Rayon can be used to create thread worker

```
fn sum_of_squares(input: &[i32]) -> i32 {  
    input.par_iter()  
        .map(|i| i * i)  
        .sum()  
}  
  
fn add_slices(a: &[f32], b: &[f32]) -> Vec<f32>{  
    let iter_a = a.par_iter();  
    let iter_b = b.par_iter();  
    a.zip(b).map(|(a, b)| a + b).collect()  
}
```

Similar to OpenMP in C. Usefull for a big number of small functions.

Async/await, Future and cooperative concurrence

Async/await, Future and cooperative concurrence

**Getting two pages at the same
time**

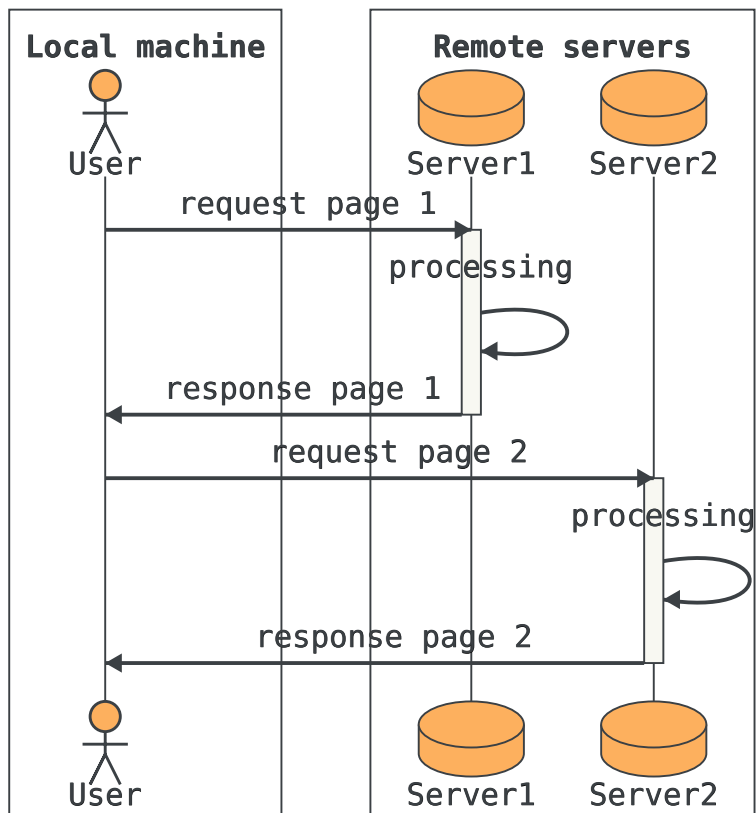
Problem statement

How to get the content of two web pages ?

- as fast as possible
- without using too much ressources

Sequential

Getting two webpage, Sequential




- simple
- slow
- use only one stack

Getting two pages at the same time

Threaded

Getting two webpage, Threaded



- use the os for context switching
- fast
- use three stacks

Issue with threads

- Compute vs IO bound:
 - here, we are IO-bound
 - adding compute time through threads is not usefull

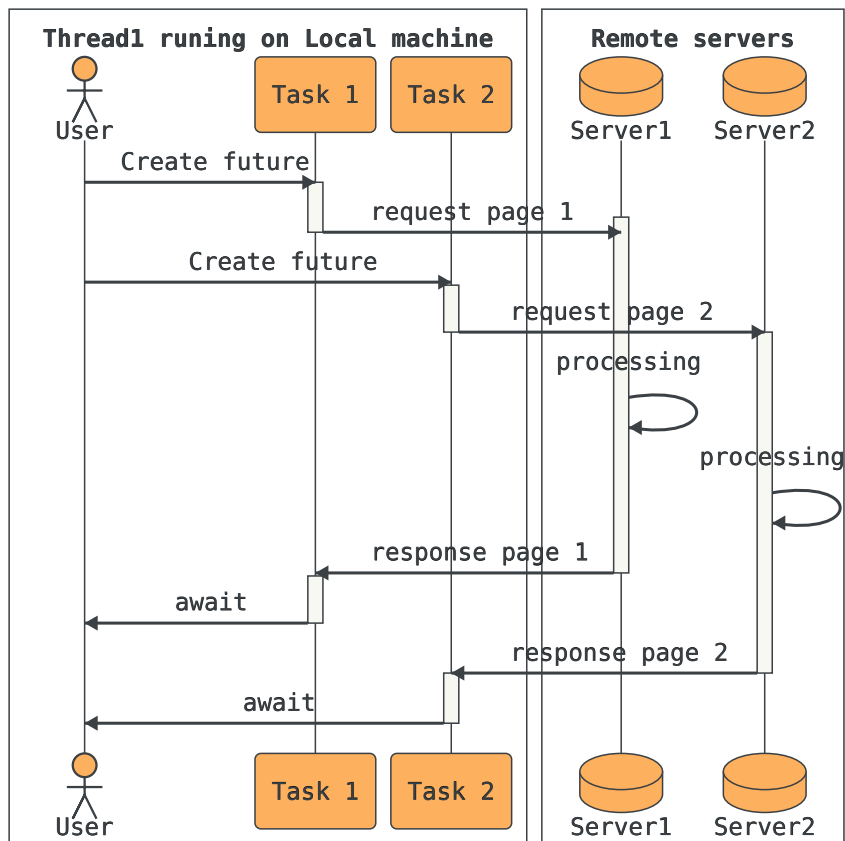
Issue with threads

- Compute vs IO bound:
 - here, we are IO-bound
 - adding compute time through threads is not usefull
- Scaling to many tasks:
 - one thread by task \Rightarrow one stack by task \Rightarrow heavy memory usage
 - could we use only one of the main thread and distribute compute time to tasks ?

Getting two pages at the same time

Asynchronous tasks

Getting two webpage, async



- use a local algorithm for context switching
- fast
- use one stack

Task, future, executor

- tasks, and their sub-tasks:
 - ▶ are represented as the promise of a future value (that value can be the unit type if there is no return value)
 - ▶ can do little progress at once, stopping when they can not progress anymore or are finished
- an algorithm, the **executor**:
 - ▶ runs on the main thread
 - ▶ distributes compute time at each task

Async/await, Future and cooperative concurrency

Implementation in Rust

```
enum Poll<T> { // core::task::Poll
    Ready(T),
    Pending,
}
trait Future { // core::future::Future
    type Output;
    // Required method
    fn poll(self: Pin<&mut Self>, cx: &mut Context<'_>)
        -> Poll<Self::Output>;
}
```

The context of a task is stored in the other member of the struct. `poll` must be **nonblocking**.

- Resolving a future is polling it until it returns Ready.
- Future implementation uses a `Waker` to signal that the future is ready to be polled.
 - this allows for the task to only wakeup when progress is possible.
- Unlike Javascript or C#, the task polling is done by the program, not by the runtime/virtual machine.

- Future should be implemented as state machine,
- Future should not start their work before the first `poll`,
- Future should use the `Waker` to be called again,
- if a Future contains another Future, it should call the `poll` of the children each time it is called,
- `poll` must be quick

```
struct Delay {when: Instant}
impl Future for Delay {
    type Output = &'static str;
    fn poll(self: Pin<&mut Self>, cx: &mut Context<'_>)-> Poll<&'static str>
    {    if Instant::now() >= self.when {
            println!("Hello world");
            Poll::Ready("done")
        } else {
            let waker = cx.waker().clone();// Get a handle to the waker
            let when = self.when;
            thread::spawn(move || { // Spawn a timer thread.
                let now = Instant::now();
                if now < when {thread::sleep(when - now);}
                waker.wake();
            });
            Poll::Pending
        }
    }
}
```

Implementation in Rust

Async function

```
// fn slow_getter(&str) -> impl Future<i32>;  
async fn slow_get_add(r: &str, n: i32) -> i32 {  
    let r = get_slow(r).await;  
    r + n  
}
```

```
// fn slow_getter(&str) -> impl Future<i32>;  
async fn slow_get_add(r: &str, n: i32) -> i32 {  
    let r = get_slow(r).await;  
    r + n  
}
```

async marks the function, the compiler will turn it into the function `fn slow_getter(&str, i32) -> impl Future<i32>`. The associated structure implementing `Future` will also be generated. The `.await` marks that the future should poll the `get_slow` future.

The function

```
async fn seek_value(name) -> Option<Elem> {  
    let iter = name.get().await;  
    for (n, elem) in iter.enumerate() {  
        if check(elem).await {return elem}  
    }  
    None  
}
```

Gives a state machine that looks like

```
enum SeekValueState{  
    StateGet(name, Get),  
    StateCheck(name, iter, n, elem, Check)  
}
```

Multiple Futures to await

```
async fn get_plus_twice_bad(r1: &str, r2: &str, n: i32) -> (i32, i32) {  
    let p1 = slow_get_add(r1, n).await;  
    let p2 = slow_get_add(r2, n).await;  
    (p1, p2)  
}  
  
async fn get_plus_twice(r1: &str, r2: &str, n: i32) -> (i32, i32) {  
    let pair = join(slow_get_add(r1, n), slow_get_add(r2, n));  
    pair.await  
}
```

Futures will only start to execute if they are awaited/pollled.

The executor is the object that handle polling the tasks.

The executor is the object that handle polling the tasks.

There is no runtime/executor in `std`. The main ones are:

- Tokio for hosted environnement
 - includes io async function for IO (network and filesystem)
 - awaitable synchronisation primitive (mutex, channel)
- Embassy for embedded environnement
 - includes async hardware abstraction layer


```
#[tokio::main]
async fn main() -> Result<()> {
    let mut client = client::connect("127.0.0.1:6379").await?;

    client.set("hello", "world".into()).await?;

    let result = client.get("hello").await?;

    println!("got value from the server; result={:?}", result);

    Ok(())
}
```

Implementation in Rust

Tokio, creating tasks

```
#[tokio::main]
async fn main() -> Result<(), Box<dyn std::error::Error>> {
    let listener = TcpListener::bind("127.0.0.1:8080").await?;
    loop {
        let (mut socket, _) = listener.accept().await?;
        let handle = tokio::spawn(async move {
            let mut buf = [0; 1024];
            loop { // In a loop, read data from the socket and write the data back.
                let n = match socket.read(&mut buf).await {
                    Ok(0) => return, // socket closed
                    Ok(n) => n,
                    Err(e) => {
                        eprintln!("failed to read from socket; err = {:?}", e);
                        return;
                    }
                };
                if let Err(e) = socket.write_all(&buf[0..n]).await {
                    eprintln!("failed to write to socket; err = {:?}", e);
                    return;
                }
            }
        });
    }
}
```

When spawned with spawn, tasks are eagerly executed.

To execute code that could block, you should use `tokio::spawn_blocking`.
This will spawn a thread and await for the end of the thread.
This allows to create async interfaces from blocking interfaces.

Tokio can use multiple threads to poll multiple futures at the same time.
For futures to be shared between worker threads, they need to be Send.

A future created with `async` is `Send` if all the variables held during an `await` are `Send`.

A future created with `async` is `Send` if all the variables held during an `await` are `Send`. Holding a mutex handle during an `await` is not `Send`.

- No blocking in async functions!

Good practice for async

- No blocking in async functions!
- Do not hold a mutex across a `await`,

- No blocking in async functions!
- Do not hold a mutex across a `await`,
- use the `join!` macro to await multiple futures at the same time.
- use the crate futures:
 - the `join` futures to await multiple futures at the same time,
 - an awaitable `Mutex`

Conclusion



Threads vs Futures

- Threads
 - Compute-bound program
 - expensive to start
 - can exploit multi-hart computer
 - relies on the OS
 - communication via channel or mutex
 - can block
- Futures
 - IO-bound program
 - cheap to start
 - might exploit multi-hart computer
 - relies on a user provided executor

Threads vs Futures

- should not block

Using the best of both worlds



- multithreaded executor :
 - ▶ use multiple threads to `poll` multiple futures at the same time,
 - ▶ it is default in tokio,

- multithreaded executor :
 - ▶ use multiple threads to `poll` multiple futures at the same time,
 - ▶ it is default in tokio,
- using a thread to resolve long calculation
 - ▶ the future creates/acquires a thread, launches a calculation and returns `pending`,
 - ▶ when the thread terminates, it wakes the future and gives the result
 - ▶ can also use a pool of threads

Questions ?